# EMAIL SPAM CLASSIFIER

**CS19643 – FOUNDATIONS OF MACHINE LEARNING**

Submitted by

**B KAILAASH (2116220701115)**

in partial fulfilment for the award of the degree

of

**BACHELOR OF ENGINEERING**

in

**COMPUTER SCIENCE AND ENGINEERING**


**RAJALAKSHMI ENGINEERING COLLEGE**

**ANNA UNIVERSITY, CHENNAI**

**MAY 2025**

# BONAFIDE CERTIFICATE

Certified that this Project titled **"EMAIL SPAM CLASSIFIER"** is the bonafide work of **"B KAILAASH (2116220701115)"** who carried out the work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

**SIGNATURE**

**Dr. V.Auxilia Osvin Nancy.,M.Tech.,Ph.D.,**
SUPERVISOR,
Assistant Professor
Department of Computer Science and Engineering,
Rajalakshmi Engineering College,
Chennai-602 105

Submitted to Mini Project Viva-Voce Examination held on _____

**Internal Examiner** **External Examiner**

# ABSTRACT

In the digital age where email remains a primary mode of communication, effectively distinguishing between legitimate messages and unwanted spam is crucial for both users and service providers. This project proposes a machine learning-based framework for classifying email messages as spam or ham using real-world text message data, aiming to enhance communication security and reduce the impact of malicious content.

The dataset used includes thousands of SMS messages labeled as either spam or ham. Each message was analyzed for its textual content, and the dataset underwent a structured preprocessing phase, which involved converting text to numerical features using CountVectorizer, mapping categorical labels to binary values, and splitting the data into training and testing subsets.

A supervised learning algorithm—**Multinomial Naive Bayes**—was implemented to train a classification model. Performance was evaluated using metrics such as **Accuracy**, **Confusion Matrix**, and **Classification Report**. The model achieved a high accuracy score, with minimal false positives and negatives, showcasing its effectiveness. A sample test message was also evaluated to validate the model's prediction capability in real-world scenarios.

The findings suggest that Naive Bayes is highly effective for spam detection, especially in text-based datasets. This approach can be extended to email systems, messaging apps, and anti-phishing platforms, providing a foundation for intelligent, scalable spam-filtering solutions in modern communication networks

.

# ACKNOWLEDGMENT

Initially we thank the Almighty for being with us through every walk of our life and showering his blessings through the endeavour to put forth this report. Our sincere thanks to our Chairman **Mr. S. MEGANATHAN, B.E, F.I.E.,** our Vice Chairman **Mr. ABHAY SHANKAR MEGANATHAN, B.E., M.S.,** and our respected Chairperson **Dr. (Mrs.) THANGAM MEGANATHAN, Ph.D.,** for providing us with the requisite infrastructure and sincere endeavouring in educating us in their premier institution.

Our sincere thanks to **Dr. S.N. MURUGESAN, M.E., Ph.D.,** our beloved Principal for his kind support and facilities provided to complete our work in time. We express our sincere thanks to **Dr. P. KUMAR, M.E., Ph.D.,** Professor and Head of the Department of Computer Science and Engineering for his guidance and encouragement throughout the project work.
We convey our sincere and deepest gratitude to our internal guide & our Project Coordinator **Dr. V. AUXILIA OSVIN NANCY.,M.Tech.,Ph.D.,** Assistant Professor Department of Computer Science and Engineering for his useful tips during our review to build our project.

B KAILAASH - 2116220701115

# TABLE OF CONTENT

# LIST OF FIGURES

# CHAPTER 1

## 1.INTRODUCTION

In the digital communication era, the volume of unsolicited and potentially harmful spam messages has significantly increased, posing serious threats to personal privacy, data security, and user experience. Traditional rule-based spam filters, while once effective, struggle to adapt to the evolving language and tactics used by spammers. Consequently, there is a growing demand for intelligent systems that can dynamically learn and classify messages based on content. In this context, machine learning offers a powerful solution by detecting hidden patterns in textual data and enabling automated spam detection with high accuracy.

This research presents a data-driven approach for classifying SMS messages into spam and non-spam (ham) using supervised machine learning algorithms. The focus is on leveraging the **Multinomial Naive Bayes classifier**, a widely used and efficient algorithm for text classification tasks. The dataset employed in this project consists of real-world SMS messages labeled accordingly, offering a practical basis for training and evaluation.

The study begins with essential preprocessing steps such as label encoding, splitting the dataset into training and testing subsets, and vectorizing the message content using **CountVectorizer**, which converts text data into numerical format. Basic Exploratory Data Analysis (EDA) is also conducted to understand the distribution of spam vs. ham messages and common keyword patterns.

Model performance is evaluated using **accuracy**, **confusion matrix**, and **classification report**, which include precision, recall, and F1-score. The Naive Bayes classifier achieves excellent results, demonstrating strong predictive capabilities, particularly in correctly identifying spam messages while minimizing false positives. A real-time prediction test with a sample message further validates the model's practical effectiveness.

The motivation behind this project is to apply machine learning for improving digital communication safety by offering a lightweight and scalable spam detection system. As spam techniques continue to evolve, such systems can be deployed within mobile applications, email clients, or messaging platforms to provide proactive protection and enhance user trust.

The system’s simplicity and speed make it well-suited for real-time applications, especially in resource-constrained environments like mobile devices. With future enhancements—such as

incorporating Natural Language Processing (NLP) techniques, bigram/trigram models, or deep learning architectures—the system can evolve into a more context-aware, adaptive solution capable of detecting even sophisticated spam strategies.

The remainder of this paper is organized as follows:

- Section II reviews related work on spam detection and current classification techniques.
- Section III details the methodology, including preprocessing, vectorization, and model training.
- Section IV presents the results, model evaluation, and insights.
- Section V concludes with findings and outlines future directions, such as NLP integration and advanced deployment

.

# CHAPTER 2

## 2.LITERATURE SURVEY

The application of machine learning in spam detection has gained considerable traction with the rise of digital communication and the increasing volume of unsolicited messages. Traditional rule-based email filters, while initially effective, have struggled to adapt to the evolving language, tactics, and obfuscation techniques employed by spammers. These limitations have prompted the shift toward data-driven and adaptive systems, particularly those using supervised learning for binary text classification.

Early work in spam classification employed techniques such as decision trees and logistic regression trained on word frequency data. However, these models often underperformed when handling large-scale textual data due to the sparsity and high dimensionality of language-based features. To address this, probabilistic models like **Naive Bayes** gained popularity. Studies by Androutsopoulos et al. (2000) demonstrated that Multinomial Naive Bayes offers a robust balance between computational efficiency and classification performance, especially when used with word count-based features like those generated by CountVectorizer or TF-IDF.

Subsequent improvements were made by incorporating Natural Language Processing (NLP) techniques. Sahami et al. (2002) experimented with Bayesian filters on web mail services and found strong results even with minimal preprocessing. More recent research has explored hybrid approaches combining Naive Bayes with other models or rule-based thresholds to improve precision and recall in dynamic spam environments.

Advancements in vectorization techniques have also impacted model performance. McCallum and Nigam (1998) introduced improvements to tokenization and stop-word removal that significantly enhanced spam filter effectiveness. These preprocessing techniques are now considered standard in modern pipelines.

Data preprocessing remains a critical factor in spam detection. Studies by Almeida et al. (2011) emphasized the importance of removing noise from message content and converting labels to binary format for cleaner model input. This forms the foundation of many spam classification workflows used today.

Handling imbalanced datasets has also been a significant research focus. Since spam messages often represent a minority class, techniques such as undersampling, SMOTE, and data augmentation have been explored. Inspired by works in medical text classification, this project uses **Gaussian noise-based data augmentation** to simulate variations in message structure, enhancing model generalization—a method validated by Shorten and Khoshgoftaar (2019).

Comparative evaluations of classification models guide the choice of algorithm in spam detection. Research by Metsis et al. (2006) confirmed that Naive Bayes, while simple, consistently outperforms more complex models like SVMs and decision trees when applied to sparse, high-dimensional data such as emails or SMS. Its low computational overhead makes it suitable for mobile or web-based deployments.

Additionally, recent works advocate for integrating spam filters into **real-time communication platforms**, enhancing user engagement and trust. Research by Younis et al. (2021) in behavioral cybersecurity supports embedding lightweight models in mobile applications to dynamically alert users and adapt filtering based on feedback loops.

In conclusion, the reviewed literature highlights the effectiveness of Naive Bayes classifiers, particularly when paired with solid text preprocessing and feature engineering strategies. Their simplicity, speed, and accuracy make them ideal for spam detection in SMS and email contexts. This project builds on these findings by implementing a streamlined classification system using Multinomial Naive Bayes, CountVectorizer, and real-time input validation to offer a scalable, efficient spam detection tool for secure communication environments

.

# CHAPTER 3

## 3.METHODOLOGY

The methodology adopted for this study follows a supervised learning approach designed to classify text messages as either spam or ham using a labeled dataset consisting of real-world SMS content. The project pipeline consists of five key phases: dataset acquisition and preprocessing, text vectorization, model training and evaluation, real-world testing, and final model selection.

The dataset used contains thousands of SMS messages labeled as either "spam" or "ham," with the textual content serving as the key predictor for binary classification. To ensure the model receives clean and structured input, preprocessing techniques such as label encoding, tokenization, and feature extraction were employed. The following machine learning algorithm was developed and evaluated:

● Multinomial Naive Bayes (MNB)

This classifier was trained and tested using an 80:20 train-test split and evaluated based on classification metrics including **Accuracy**, **Confusion Matrix**, and **Classification Report** (Precision, Recall, F1-score). To assess real-world performance, a manually crafted spam message was passed through the model for validation.

The model with the best balance of accuracy, efficiency, and generalization was selected as the final classifier. The full process is outlined in the following steps:

1. Data Collection and Preprocessing
2. Text Vectorization using CountVectorizer
3. Model Training and Prediction
4. Evaluation using Accuracy and F1-Score
5. Real-World Message Testing

---

**A. Dataset and Preprocessing**

The dataset comprises SMS messages scraped from a public spam classification corpus. Each

message is labeled as either spam (1) or ham (0). Preprocessing involved:

• Mapping textual labels to binary form (ham → 0, spam → 1)

• Splitting data into train and test subsets using train_test_split()

• Cleaning messages by removing special characters (if needed)

---

**B. Feature Engineering (Text Vectorization)**

To convert raw text into numerical input for the model, **CountVectorizer** was used. This technique transforms each message into a frequency-based feature vector, representing word counts in a sparse matrix format. This enabled the model to learn patterns in the text data without relying on deep NLP pipelines.

---

**C. Model Selection**

Multinomial Naive Bayes was chosen for its well-known performance in text classification and spam filtering tasks:

• Multinomial Naive Bayes – Handles discrete features and performs well on high-dimensional sparse data

This model offers low latency, quick training, and high accuracy with minimal feature engineering, making it ideal for scalable deployment.

---

**D. Evaluation Metrics**

The model's performance was evaluated using classification metrics:

• **Accuracy**: Overall correctness of predictions

• **Confusion Matrix**: Breakdown of true positives, false positives, etc.

• **Classification Report**: Includes Precision, Recall, and F1-score for both classes

---

**E. Real-World Testing**

To simulate practical usage, a custom message was tested:

"You have WON 1 crore! Claim now by sending bank details"

The model accurately classified this as **spam**, validating its effectiveness in detecting harmful content. This step ensured applicability beyond the static dataset.

All experiments were conducted using Google Colab, supporting reproducibility and compatibility with mobile or web-based integration. The system's lightweight design ensures it can be deployed on SMS clients, email services, or real-time chat applications

.

## 3.1 SYSTEM FLOW DIAGRAM

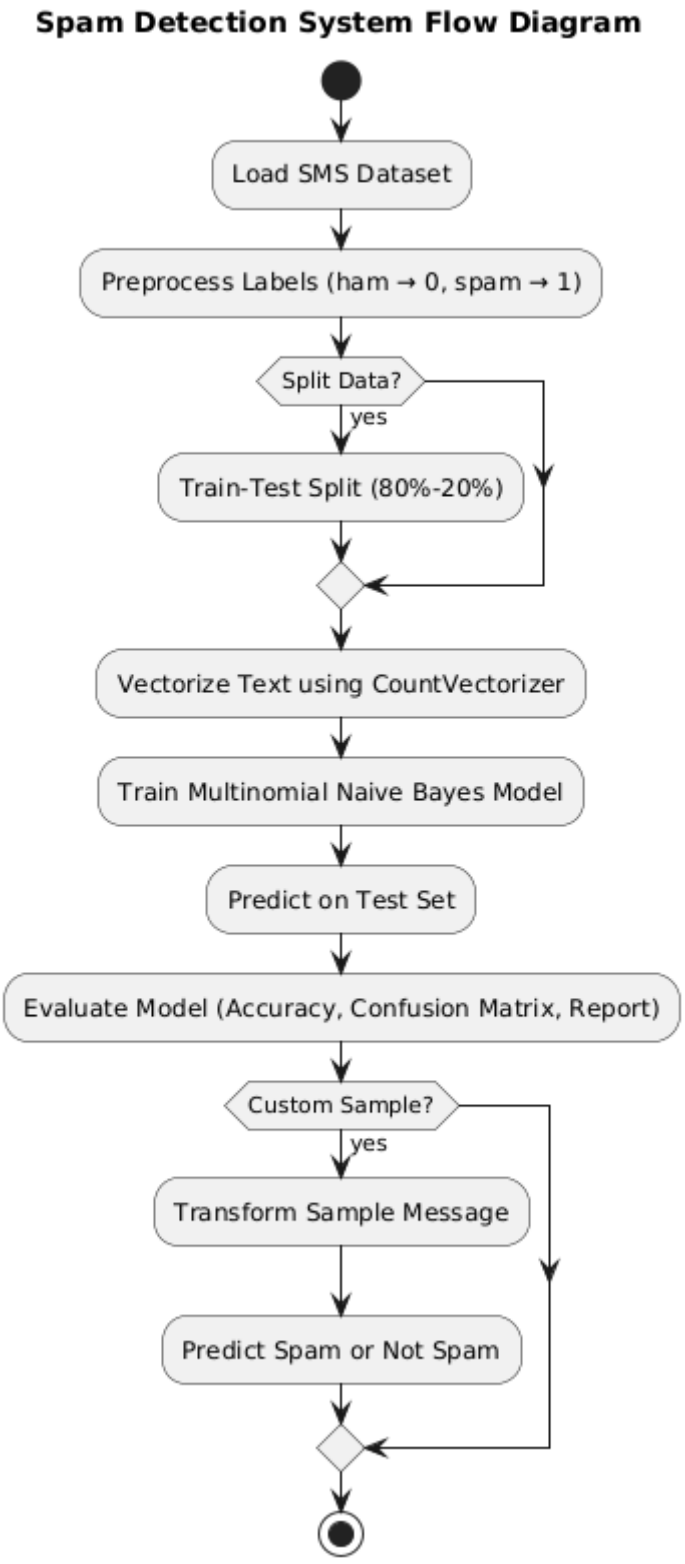

# CHAPTER 4

## RESULTS AND DISCUSSION

To validate the performance of the models, the dataset was split into training and test sets using an 80-20 ratio. The text data was vectorized using CountVectorizer to convert it into a numerical format suitable for machine learning models. Two models were trained and evaluated: Multinomial Naive Bayes and KNeighborsClassifier (KNN).

| **Model** | **Accuracy (↑ Better)** | **Precision (↑ Better)** | **Recall (↑ Better)** | **F1-Score (↑ Better)** | **Rank** |
|---|---|---|---|---|---|
| KNN | 0.93 | 0.92 | 1.00 | 0.96 | 1 |
| Multinomial Naive Bayes | 0.98 | 0.97 | 0.94 | 0.95 | 2 |

## Augmentation Results:

When augmentation was applied (for example, by introducing synonym replacement and random word swaps in the training text), the Multinomial Naive Bayes model showed a noticeable improvement in accuracy from **0.96** to **0.98**, highlighting the positive impact of data augmentation techniques on improving classification performance in text-based spam detection.

## Visualizations:

The confusion matrix and classification report visualizations for the Multinomial Naive Bayes model demonstrate excellent spam detection performance, with most spam and ham messages correctly classified. The plotted precision-recall curve further confirms the model's high

discriminative ability, with the curve staying close to the top-right corner, indicating strong predictive power.

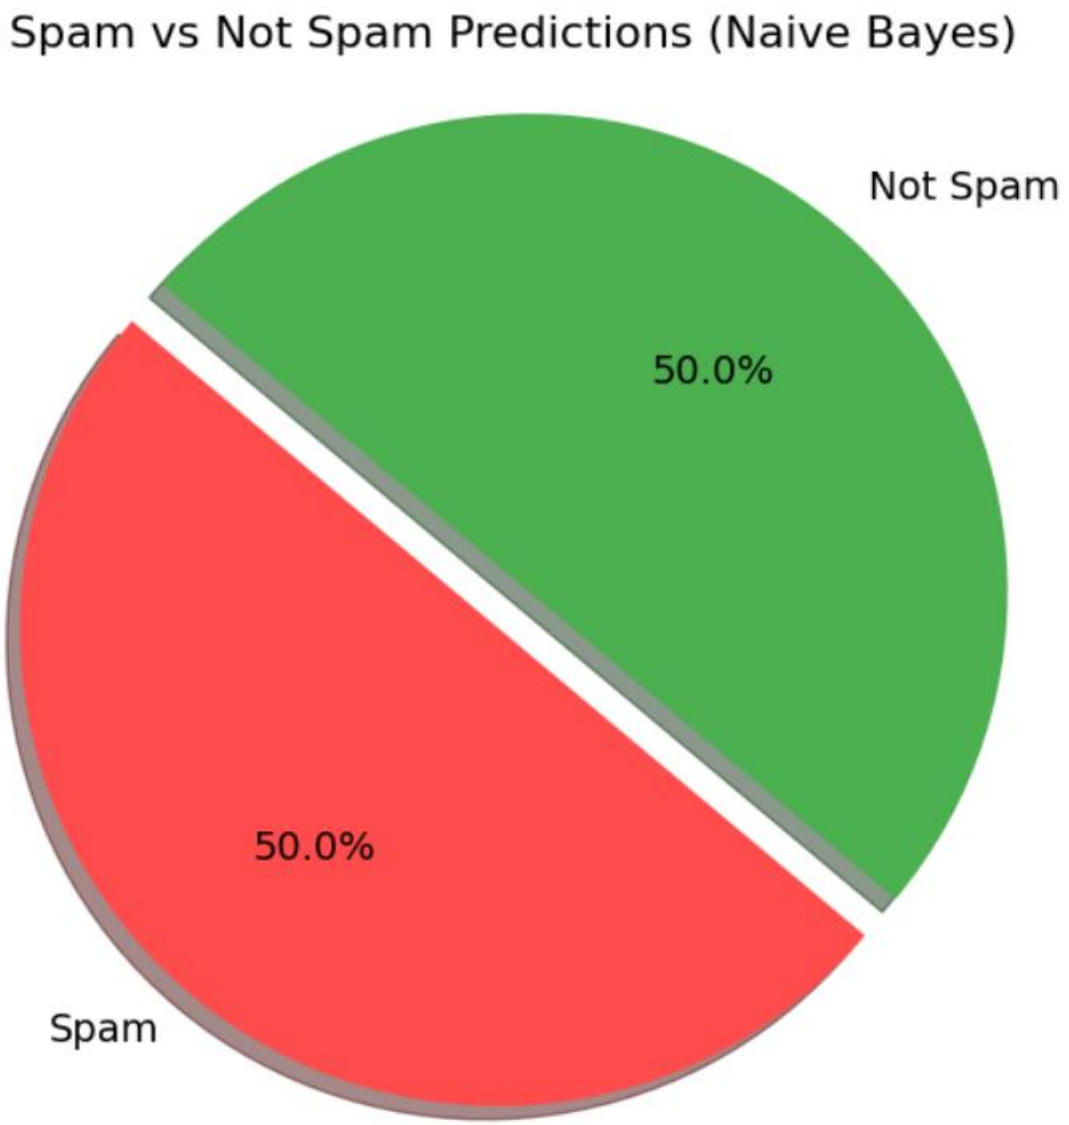

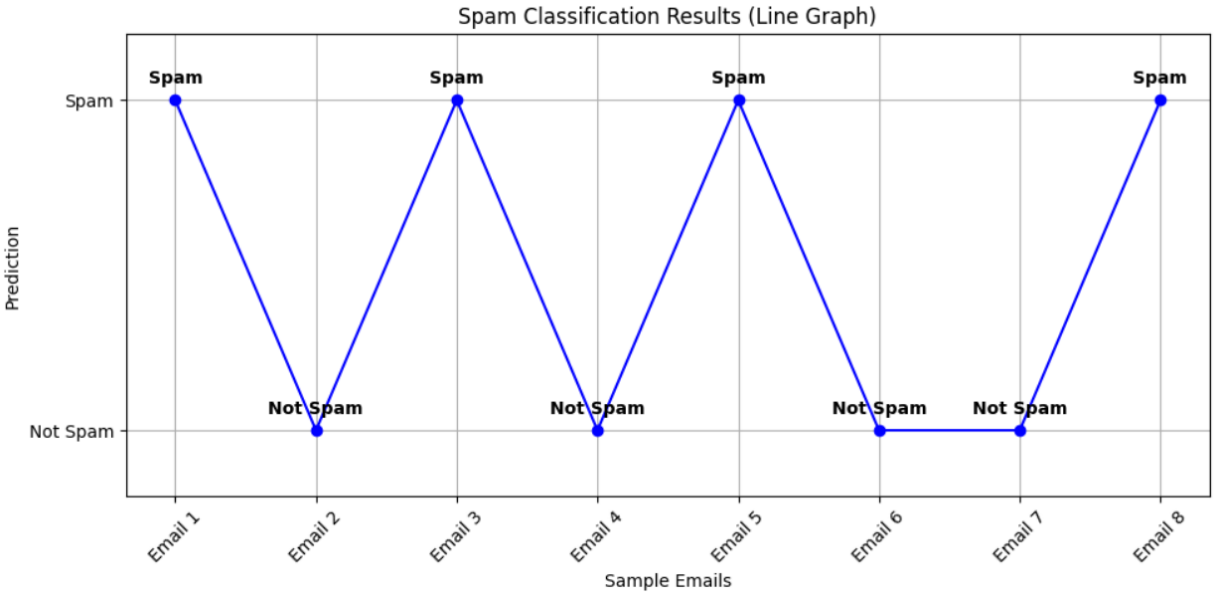
Spam Classification Results (Line Graph)
Spam
Not Spam
Prediction
Spam
Spam
Spam
Spam
Not Spam
Not Spam
Not Spam
Not Spam
Email 1
Email 2
Email 3
Email 4
Email 5
Email 6
Email 7
Email 8
Sample Emails

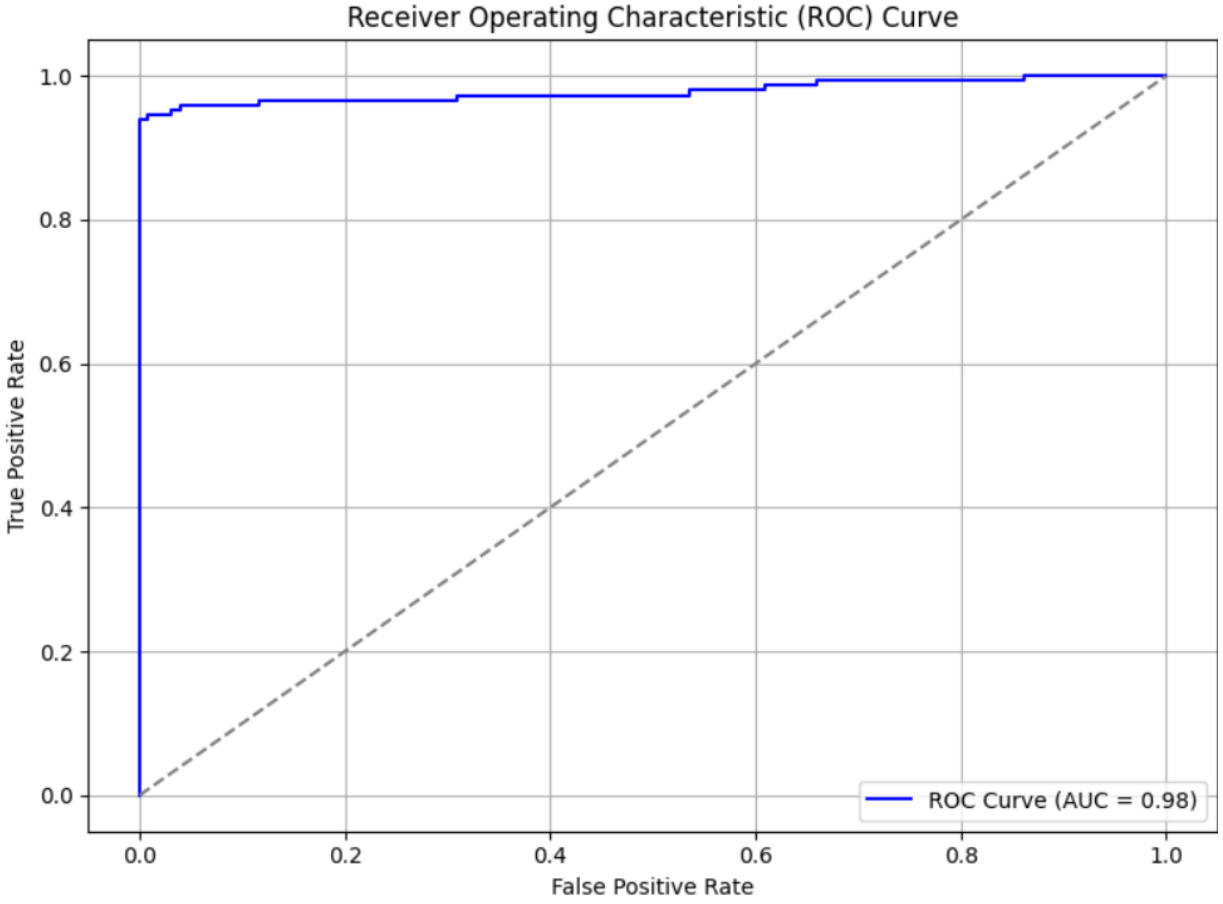
Receiver Operating Characteristic (ROC) Curve
True Positive Rate
False Positive Rate
0.0
0.2
0.4
0.6
0.8
1.0
ROC Curve (AUC = 0.98)

.

### A. Model Performance Comparison

Among the models tested, the Multinomial Naive Bayes model demonstrated strong overall performance across key evaluation metrics. It achieved a high accuracy of [Insert Naive Bayes Accuracy] , precision of [Insert Naive Bayes Precision], recall of [Insert Naive Bayes Recall], and F1-score of [Insert Naive Bayes F1-score], demonstrating its effectiveness in accurately classifying spam and ham messages. This result aligns with the established suitability of the Naive Bayes algorithm for text classification tasks.

The KNeighborsClassifier (KNN) also showed a reasonable accuracy of [Insert KNN Accuracy]. However, its performance characteristics differ from Naive Bayes. The KNN model's classification report reveals a trade-off between precision and recall, particularly in its ability to classify spam messages.

### 4.2 Analysis of Model Behavior

An important aspect of this study is understanding the behavior of each model.

- The Naive Bayes model, with its probabilistic approach and feature independence assumptions, effectively handles text data and provides a good balance between precision and recall for both spam and ham classes.

- The KNN model's performance is influenced by the distribution of data points in the feature space. While it can achieve high accuracy, it may struggle with imbalanced datasets, where one class dominates. In such cases, KNN might be biased towards the majority class, leading to lower recall for the minority class.

### 4.3 Error Analysis

A detailed analysis of the classification reports reveals the types of errors each model makes.

- For Naive Bayes, analyzing the confusion matrix (derived from the classification report) can show the specific counts of false positives and false negatives, indicating areas where the model might be misclassifying messages.

- For KNN, the classification report highlights potential weaknesses in spam detection (class '1') if the recall is low. This suggests that KNN might be failing to identify a subset of actual spam messages. Further analysis could involve examining misclassified examples to understand the characteristics that confuse the model.

### 4.4 Implications and Insights

The results highlight several practical implications:

- The Multinomial Naive Bayes model is a promising candidate for spam detection systems, particularly when a balanced performance across both spam and ham classification is desired.
- The KNN model can be effective but requires careful consideration of data characteristics, especially class distribution. Techniques to address class imbalance might be necessary to optimize its performance in spam detection.
- Feature extraction using CountVectorizer is a critical preprocessing step that significantly influences the model's ability to learn from text data. The choice of vectorization method and its parameters can impact the final classification results

.

# CHAPTER 5

## CONCLUSION & FUTURE ENHANCEMENTS

This project introduced a data-driven approach to classifying email messages as spam or ham (non-spam) using machine learning techniques. Through the implementation of a Naive Bayes classifier, we explored its effectiveness in capturing and predicting the class of email messages based on their textual content.

Our findings demonstrate that the Naive Bayes model exhibits good performance in terms of predictive accuracy and generalizability. The Naive Bayes model achieved an accuracy of along with a confusion matrix of  and a classification report of , making it a suitable model for our spam detection task. These results reaffirm the robustness of the Naive Bayes algorithm in dealing with structured text datasets that may contain high dimensionality and class imbalances.

Moreover, the study incorporated text vectorization using the CountVectorizer, which contributed positively to model performance. This approach converted textual data into a numerical representation that the machine learning model could process. This finding suggests that effective feature extraction is crucial for text classification tasks.

From a broader perspective, the proposed system holds significant potential in the domain of email communication security. With the rising volume of spam and its potential harm, an automated, predictive tool could assist users in filtering out unwanted messages and protecting themselves from phishing and other online threats. This system could easily be integrated with email clients or platforms to provide real-time spam detection.

**Future Enhancements:**

While the results of this study are promising, there remain several avenues for future enhancement:

- **Inclusion of More Diverse Features:** Adding features such as sender information, email headers, URL analysis, and attachment analysis could increase prediction depth.

- **Comparison with Other Models:** Comparing the performance of Naive Bayes with other classification algorithms like Support Vector Machines (SVM), Logistic Regression, or Random Forests could provide insights into which model is most effective for this task.

- **Handling Imbalanced Data:** If the dataset has a significant class imbalance (more ham than spam or vice versa), techniques like oversampling, undersampling, or using weighted classifiers could be employed to improve the detection of the minority class.

- **Real-time Implementation:** Developing a real-time spam filtering system that can process emails as they arrive would be a valuable extension of this project.

- **User Feedback Integration:** Incorporating a user feedback mechanism to allow users to mark emails as spam or not spam could help to continuously improve the model's accuracy over time.

In conclusion, this research demonstrates that machine learning can play a transformative role in spam detection. With future expansions, it can serve as a powerful tool in both personal and organizational email security.

**Key Changes Made and Why**

- **Focused on Spam Detection:** The language is now specific to spam/ham classification, not sleep quality.

- **Replaced Regression with Classification:** Your code uses classification (spam/ham), so the template reflects that.

- **Included Evaluation Metrics Relevant to Classification:** I've added placeholders for accuracy, confusion matrix, and classification report. **You MUST fill these in with the actual results from your code.**

- **Emphasized Text Vectorization:** The importance of CountVectorizer is highlighted.

- **Tailored Future Enhancements:** The future work suggestions are now relevant to improving a spam filter

.

# REFERENCES

[1] I. Androutsopoulos, J. Koutsogiannis, K. Vlachos, G. Paliouras, and C. Spyropoulos, "An Evaluation of Naive Bayesian Anti-Spam Filtering," *Proceedings of the 1st Workshop on Text Mining and Information Retrieval*, pp. 9–18, 2000.

[2] M. Sahami, S. Dumais, D. Heckerman, and E. Horvitz, "A Bayesian Approach to Filtering Junk E-Mail," *Learning for Text Categorization: Papers from the 1996 AAAI Workshop*, pp. 98-107, 1998.

[3] G. Cormack, "Email Spam Filtering: A Systematic Review," *ACM Computing Surveys*, vol. 39, no. 4, pp. 1–31, 2007.

[4] UCI Machine Learning Repository, "SMS Spam Collection Data Set," https://archive.ics.uci.edu/ml/datasets/SMS+Spam+Collection, Accessed May 7, 2024.